

Correct-by-Construction Reinforcement Learning of Cardiac Pacemakers from Duration Calculus Requirements*

Kalyani Dole^{†1}, Ashutosh Gupta^{†1}, John Komp^{†2},
Shankaranarayanan Krishna^{†1}, Ashutosh Trivedi^{†2}

¹ Department of Computer Science and Engineering, Indian Institute of Technology Bombay

Email: {kalyanid,akg,krishnas}@cse.iitb.ac.in

² Department of Computer Science, University of Colorado Boulder

Email: {john.komp,ashutosh.trivedi}@colorado.edu

Abstract

As the complexity of artificial pacemakers continues to grow, the importance of capturing its functional correctness requirement formally cannot be overestimated. The pacemaker system specification document by *Boston Scientific* provides a widely accepted set of specifications for pacemakers. As these specifications are written in a natural language, they are not amenable to automated verification, synthesis, or reinforcement learning of pacemaker systems. This paper presents a formalization of these requirements for a dual-chamber pacemaker in *duration calculus* (DC), a highly expressive real-time specification language. The proposed formalization allows us to automatically translate pacemaker requirements into executable specifications as stopwatch automata, which can be used to enable simulation, monitoring, validation, verification and automatic synthesis of pacemaker systems. The cyclic nature of the pacemaker-heart closed-loop system results in DC requirements that compile to a decidable subclass of stopwatch automata. We present shield reinforcement learning (shield RL), a shield synthesis based reinforcement learning algorithm, by automatically constructing safety envelopes from DC specifications.

1 Introduction

The human heart is arguably the most important real-time safety-critical system. In an average life-span of 70 years, the human heart beats more than 2.5 billion times. Each heartbeat is a complex chemical reaction that starts at the *sinoatrial node* in the right atrium and sweeps down through the ventricles. The cells of the myocardium have unique properties that regulate the speed of this chemical reaction to synchronize each chamber's beat with the hemodynamics of the cardiac cycle (Kay and Shepard 2017). Any disruption in the cellular chain can cause the heart to beat improperly.

*This material is based upon work supported by the National Science Foundation (NSF) under grant No. CCF-2009022 and NSF CAREER award CCF-2146563.

[†]These authors contributed equally.

Copyright © 2023, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

A common cardiac arrhythmia, known as heart block or AV block, is a condition where the electrical signal controlling the ventricular contractions are partially or completely blocked (Vijayaraman and Ellenbogen 2017). Depending upon the severity of this block, pacemaker therapy may be indicated. A pacemaker is an implantable electronic device that sends electrical impulses to the heart to regulate the heartbeat. The correctness of the pacemaker-heart closed-loop system is time-critical: it should not only provide supportive pacing when needed, it must ensure that the heart is not sent into unnaturally fast heart beats. Furthermore, there are secondary requirements to improve the quality of life such as detecting and supplying increased rates when the patient exercises (rate-adaptive pacing) and maximizing battery longevity (2.2% risk of infection; 0.4% mortality rate per pacemaker replacement (Polyzos, Konstantelias, and Falagas 2015)).

Over sixty years of research and development in designing artificial cardiac pacemakers has resulted in remarkable robustness, convenience, and acceptance of these devices. As the number of patients using pacemakers continues to rise (Greenspon et al. 2012), so are the expectations on the invisibility of its design resulting in ever-increasing demands on improved functionality, size shrinkage, power consumption, battery performance and adaptability to individual cardiac arrhythmias (Gomes 2020; Finkelmeier 1991). The medical device industry continues to meet these challenges by extending its design with more sensors (e.g., for rate augmentation, sleep apnea, and haemodynamic status) (Lau, Siu, and Tse 2017) resulting in ever-increasing complexity of the device software.

While manual design, supported by exhaustive model-based development, has served well for previous generations of devices, the need for adaptability without sacrificing the basic functional correctness is ever-present (Maisel et al. 2001). This paper proposes a *correct-by-construction* reinforcement learning (RL) paradigm to design cardiac pacemakers by integrating formal requirements in RL.

Example 1 (The need for Adaptive Pacemaker). *Mobitz II*

second-degree AV block (Langendorf and Pick 1968) is a condition where the atrioventricular node does not transmit all atrial depolarizations to the ventricle and is manifested as a set pattern of dropped beats. A 2:1 heart block drops every other beat while 3:2 block drops every third. While the standard pacing algorithms fill in these missing beats providing patient support, the timing of the needed ventricular paces are governed by static decisions (pace at the lower rate interval) aimed towards satisfying the safety requirements and do not attempt to match the intrinsic heart rate. As the patients can feel this difference in heartbeats, their quality of life can be enhanced by adapting the time to provide for a missing beat to the patient intrinsic rhythm.

The overarching objective of this industrial collaboration is to formalize the requirements in a rigorous, formal language capable of precisely capturing the requirements of a cardiac pacemaker. The key requirements on a pacemaker include maintaining a minimum heart rate, avoiding pacing during the ventricular repolarization, not increasing the heart rate excessively, and not disturbing the natural rhythm of the heart. The precise algorithms to implement these requirements require robust detection of various events as well as a careful accounting of a number of timers to provide the control logic. *Pacemaker System Specification* document (Laboratory 2007) by Boston Scientific provides a refined version of the basic specifications of the pacemaker design in a natural language. A key contribution of this paper is capturing these time-critical requirements for the dual-chamber pacemaker in a real-time formal logic.

Previous Efforts. Research into using pacemakers as a vehicle for applying formal methods to medical device design traces back to the *Pacemaker Grand Challenge* in 2007. The challenge was kicked off with the release of a generic dual chamber pacemaker specification (Laboratory 2007) from Boston Scientific detailing the functional correctness requirements on the system. Before release, Boston Scientific removed references to proprietary algorithms leaving only the high level intention of such features. *The Pacemaker Challenge: Developing Certifiable Medical Devices* was presented as a Schloss Dagstuhl seminar in 2014 (Méry, Schätz, and Wasssyng 2014). In 2018, Bonfanti et al. (Bonfanti, Gargantini, and Mashkoor 2018) reviewed the usage of formal methods in medical devices. Table 4 of (Bonfanti, Gargantini, and Mashkoor 2018) provides a breakdown of 48 papers related to pacemaker research into categories of modeling, model verification and validation, software validation, and code generation. A number of efforts have focused on capturing these specifications using formal languages or extensions of those languages such as AADL (Larson 2014) and Z (Gomes and Oliveira 2009). There have been attempts to model a dual-chamber pacemaker with advanced features using timed automata (Jiang et al. 2012). Timed automata are useful in capturing key features of closed-loop systems and enable the use of tools like UPPAAL in verification. While timed automata based representations are amenable to formal analysis, the translation to such specifications is manual and in the process lose interpretability; moreover, timed automata are not expressive

enough to capture several properties of interest. On the other hand, expressive specification languages limit automation.

Our Choice of Logic. We propose the use of *Duration Calculus* (DC) (Chaochen, Hansen, and Sestoft 1993; Chaochen, Hoare, and Ravn 1991) in expressing hard real-time constraints on pacemakers. DC is a real-time logic designed to express complex, time-critical properties of hybrid dynamical systems (Chaochen, Hoare, and Ravn 1991; Dole, Gupta, and Krishna 2020). The duration modalities $\int P \bowtie c$ as well as $\ell \bowtie c$ in DC allows one to capture the accumulated duration of time when some event P has been holding, as well the real time duration ℓ of interest. These modalities make DC very expressive. The expressiveness of DC, on the negative side, contributes to undecidability of key decision procedures such as emptiness and model checking against real-time models such as timed automata (Chaochen, Hansen, and Sestoft 1993); on the other hand, there has been work in exploring decidable sub classes of DC (Pandya 2002; Dole, Gupta, and Krishna 2020). There are also tools, such as DCVALID (Pandya 2001), to compile discrete DC specifications to automatic structures. To the best of our knowledge, most timed logics (other than DC) cannot model real time properties having durations. DC naturally has this modality. Circumventing the undecidability of DC by limiting the expressiveness which still is good enough to capture specifications of safety critical systems (such as the pacemaker) is a key challenge in developing correct-by-construction synthesis.

Correct-by-Construction Synthesis via RL. Correct-by-construction synthesis (Baier and Katoen 2008) is an approach to safety-critical system design that advocates the integration of the formal proof-of-correctness of the designed system by automatically refining formal specifications. It takes an adversarial view of the environment and employs tools from competitive game theory to design theoretically optimal, if over-cautious, systems. Reinforcement learning (Sutton and Barto 2018) (RL) paradigm offers an alternative view of the environment as a stochastic player with unknown strategy and proposes an adaptive sampling-based approach to converge towards the optimal policy. While RL adapts to the changes in the environment, it requires strong assumption on the nature of the environment (Markovian assumption) and lacks guarantees on the safety of the system during learning. Shielding (Alshiekh et al. 2018; Könighofer et al. 2020) is an approach to RL that combines guarantees from the correct-by-construction synthesis with adaptive nature of reinforcement learning. By focusing on cardiac pacemaker as our case-study, we develop shielding-based RL for a subclass of DC specifications.

Contributions. The contributions of this paper are summarized below.

1. We express DDD pacemaker requirements from (Laboratory 2007) in Duration Calculus.
2. We observe that due to the cyclic nature of the pacemaker-heart closed-loop system, the aforementioned requirements belong to bounded-time fragment of DC. We show that for bounded-time DC, the formulae can

be compiled into a decidable class of (acyclic) stopwatch automata (Cassez and Larsen 2000), providing a set of *executable specification*. We show that the emptiness or acyclic stopwatch automata is NP-COMplete.

3. We implemented these translations using an extension of DCVALID (dcvalid) that compiles DC specifications into stopwatch, timed, and finite-state automata.
4. We validate the behavior of the executable specifications (automata) obtained from the DC requirements for the pacemaker by showing their performance over two challenging scenarios.
5. We provide a proof-of-concept in using formal pacemaker requirements in DC to generate an RL-shield (Alshiekh et al. 2018) to restrict the behavior of the RL agent to enable safe learning. We sketch the design of an adaptive pacemaker via RL that adapts its behavior in accordance to the natural rhythm of the patient.

All of the relevant artifacts (tool and benchmarks) can be found at XXXXX.

2 Preliminaries

Duration Calculus. Duration Calculus (Chaochen, Hansen, and Sestoft 1993) is a highly expressive and succinct logic, which can capture specifications involving durations. The chop modality (\frown) contributes to its succinctness and enables compositional specification, while the measurement constructs ($\ell \bowtie c$) and ($\int P \bowtie c$) enable duration measurements.

Definition 1 (DC: Syntax). *Given a set Var of real time signals, we define the syntax of DC formulae as follows:*

$$\begin{aligned} P &::= \text{false} \mid \text{true} \mid x \in Var \mid P \wedge P \mid \neg P \\ D &::= [P] \mid [P]^\bullet \mid D \wedge D \mid \neg D \mid D \frown D \mid M \\ M &::= \ell \bowtie c \mid \int P \bowtie c \mid \sum P \bowtie c \end{aligned}$$

where $\bowtie \in \{<, \leq, =, \geq, >\}$.

Definition 2 (DC: Semantics). *Let σ be a timed trace $\langle (s_0, \tau_0), (s_1, \tau_1), \dots, (s_n, \tau_n) \rangle$ where each s_i is a state (a set of propositional variables Var) and τ_i is its time stamp. For the timed trace σ and a propositional logic formula P over Var , we say $(\sigma, i) \models P$ iff $s_i \models P$. A timed trace σ satisfies a DC formula ψ in an interval $I = [b, e]$ (where $b \leq e$ and $b, e \in \mathbb{N}$ range over the indices of a timed trace), and we write $(\sigma, [b, e]) \models \psi$, if:*

- $(\sigma, [b, e]) \models [P]$ iff $b < e$, and $(\sigma, t) \models P$ for all $b < t < e$;
- $(\sigma, [b, e]) \models [P]^\bullet$ iff $b = e$ and $(\sigma, b) \models P$;
- $(\sigma, [b, e]) \models D_1 \wedge D_2$ iff $(\sigma, [b, e]) \models D_1, D_2$;
- $(\sigma, [b, e]) \models \neg D$ iff $(\sigma, [b, e]) \not\models D$;
- $(\sigma, [b, e]) \models D_1 \frown D_2$ iff there is $b \leq z \leq e$ s.t. $(\sigma, [b, z]) \models D_1$ and $(\sigma, [z, e]) \models D_2$;
- $(\sigma, [b, e]) \models \ell \bowtie c$ iff $(\tau_e - \tau_b) \bowtie c$ holds;
- $(\sigma, [b, e]) \models \int P \bowtie c$ iff $\sum \{ \tau_{i+1} - \tau_i \mid (\sigma, i) \models P \} \bowtie c$;
- $(\sigma, [b, e]) \models \sum P \bowtie c$ iff $\{ i : (\sigma, i) \models P \} \mid \bowtie c$;

One can derive operators \Rightarrow (conditional) and \Leftrightarrow (biconditional) in the usual manner. Moreover, the temporal logic modalities eventually $\diamond D \stackrel{\text{def}}{=} \text{true} \frown D \frown \text{true}$ and the globally $\square D \stackrel{\text{def}}{=} \neg \diamond \neg D$ can be derived from the basic syntax.

Stopwatch Automata. Alur and Dill (Alur and Dill 1994) generalized the theory of finite state automata to model time-constrained evolution of systems. The resulting formalism, known as timed automata, express time constraints by using a finite set of non-negative real-valued variables called *clocks* that work as timers in that they grow with uniform rate and can be reset to 0 to remember the time since some given event. These clocks can be used in *guard expressions* on the transitions using the following grammar:

$$\varphi := x \bowtie c \mid x - x' \bowtie c \mid \varphi \wedge \varphi \quad (1)$$

where $\bowtie \in \{\leq, <, =, >, \geq\}$, x and x' are clock variables and c is a natural number. Let $\mathcal{G}(X)$ be the set of guard expressions over the set of clocks X . The stopwatch automata (Cassez and Larsen 2000) generalize timed automata by allowing the clocks to be paused (it is customary to refer to pausable clocks as stopwatches); however, adding this expressiveness results in undecidability of the emptiness problem. See Figure 2 for a stopwatch automata having a clock y and a stopwatch variable x_p .

Definition 3 (Stopwatch Automata). *A stopwatch automaton is a tuple $\mathcal{T} = (Q, q_0, \Sigma, X, \gamma, E, \delta)$ where: Q is the finite set of locations; $q_0 \in Q$ is the initial location; Σ is a finite alphabet of actions; X is the finite set of clock variables; $\gamma : Q \rightarrow 2^X$ is the set of paused clocks per location; $E : Q \times \Sigma \rightarrow \mathcal{G}(X)$ is the action guard, and $\delta : Q \times \Sigma \rightarrow \mathcal{D}(2^X \times Q)$ is the transition function.*

The DC specifications can be compiled into an “executable specification” of stopwatch automata. We omit the details of this translation due to a lack of space.

Theorem 2. *For every DC formula ϕ , one can effectively construct a stopwatch automaton \mathcal{A}_ϕ that accept the same set of timed traces.*

On a negative side, the satisfiability of DC (Chaochen, Hansen, and Sestoft 1993) as well as its emptiness (Hen-zinger et al. 1995) of stopwatch automata is undecidable.

3 Bradycardia Pacing in DC

The goal of cardiac pacing is to replace a biological component of the cardiac cycle that has failed, is operating intermittently, or to mitigate an incorrect conduction pathway that is causing irregular heartbeats. To understand how a pacemaker works starts with an understanding of the cardiac cycle. We provide a brief overview of the hear conduction system (Kay and Shepard 2017).

A heart beat begins at the sinoatrial node (SA node), causing a depolarization of cells in atrium, which compresses the atrium and forces blood into the ventricle. This phase is depicted as P-wave in Figure 1. The atrioventricular node (AV node) then delays the passage of repolarization signal to ventricle. This prevents the premature ventricular contraction (PVC) and allows blood to fill the ventricle. Then, the depolarization continues down the ventricle, causing it to compress and push the blood through the body. On the EGM, this depolarization is seen as the QRS complex. Repolarization of the ventricular cells produces a larger signal in surface EGM and is called a T-wave.




Figure 1: EGM: Cardiac Segments

Irregular heartbeats caused by a malfunctioning of SA or AV are known as bradycardia. A malfunction of the SA can lead to conditions where the heart beats slowly, irregularly, drops beats, beats rapidly, or pauses. When the AV node fails to pass the signal through the ventricle, it can result in heart block, resulting in dropped beats to complete blockage of all signals. Finally, myocardial infarction (heart attacks) can lead to the death of heart cells along the conduction path.

Modern pacemakers have mechanisms for monitoring heart intrinsic activity, transmitting generated pacing pulses and sensing the synchronization on intrinsic heart activity. A pacemaker that adds sensing and synchronization in the ventricle has the code VVI (for pace in the ventricle (V), sense intrinsic cardiac activity in the ventricle (V), and inhibit (I) a scheduled pace if a valid intrinsic ventricular event occurs). Dual chamber pacemakers (NASPE (Bernstein et al. 2002) with code DDD) pace in both the atrium and ventricle (D), sense in both chambers (D), and individual paces may be inhibited or triggered based on intrinsic activity (D). In standard pacing therapy, it is desired to pace only for the minimum required and rely on the patient's natural heart rhythm when present. This is both physiologically better for the patient and conserves the pacemaker battery for better longevity. In the absence of intrinsic activity, the pacemaker outputs a pulse to the relevant chamber at a specified rate.

At the core of pacemaker lies its ability to detect improper activity by comparing sensed intrinsic activity with the status of various period and interval timers. This helps it to decide when intrinsic events should be accepted, or to generate a pace in the future to maintain proper heart rhythm.

Due to the cyclic nature of the heart and pacemaker interactions (every atrial pace or sense restarts the timing obligations and each duration is bounded by an upper-rate interval), we observe that the pacemaker requirements fall into time-bounded fragment of DC. Motivated by the general principle of time bounded verification, (Ouaknine, Rabbinovich, and Worrell 2009), and the nature of specification needed for the DDD pacemaker, we study duration calculus whose models are evaluated over bounded timed traces.

4 Bounded Semantics for Duration Calculus

In this section, we focus on a subclass of DC whose models are evaluated over bounded timed traces. That is, DC formulae are evaluated over words of the form $\sigma = (s_0, \tau_0)(s_1, \tau_1) \dots (s_n, \tau_n)$ for some fixed and bounded n and an interval $[b, e]$. Apart from this, the DC semantics is

as before. We refer to this semantics *bounded semantics*.

A key step in developing correct-by-construction learning (and synthesis) for this subclass is its reduction to an executable specification (stopwatch automata). From this reduction, we derive the decidability of the satisfiability (reachability) and controller synthesis (safety game) problems.

4.1 From DC to Stopwatch Automata

For a DC formula D under the bounded semantics for a bound n , we construct a timed or stopwatch automaton $\mathcal{A}_D^{[0,n]}$ which accepts all timed traces of length n satisfying D . This is done inductively by building automata $\mathcal{A}_\varphi^{[i,j]}$ for $0 \leq i \leq j \leq n$ and subformulae φ of D . These automata accept timed traces of length $j-i$ satisfying φ . Depending on the subformula φ , $\mathcal{A}_\varphi^{[i,j]}$ is either a timed or a stopwatch automata. For instance if φ has the duration construct $\int P \bowtie c$, then we require stopwatches to measure accumulated durations, and otherwise, a timed automata suffices. In all cases, $\mathcal{A}_\varphi^{[i,j]}$ is an acyclic since it accepts behaviours of bounded length. The automaton construction is inductive, and we sketch key steps next.

Let Var be the set of propositional variables in D . For a Boolean combination P of variables in Var , let Var_P denote all subsets of Var which satisfy P . For example, if $Var = \{p, q, r\}$, and $P = \neg p \wedge q \Rightarrow r$, then $\{q, r\}, \{p\}, \{r\}, \{p, r\}, \emptyset \in Var_P$. We illustrate the construction of $\mathcal{A}_D^{[i,j]}$ for the cases when D is one of $\int P \bowtie c$ or $D_1 \frown D_2$. The other cases are omitted for lack of space.

1. Let $D = \int P \bowtie c$. Here, $\mathcal{A}_D^{[i,j]}$ is a stopwatch automata. Let x_P be the stopwatch variable used. From each location, we have two transitions : one on Var_P and another, on $Var_{\neg P}$. If we have a transition decorated with Var_P , then the rate of x_P is set to 1 in the target location and the rate of x_P is set to 0 in the target otherwise. It is easy to see that x_P accumulates the real time duration of P being true. After $j-i$ transitions along a path, we take the last transition to the final location, and check for $x_P \bowtie c$. If $(x_P \bowtie c)$ does not hold, we do not reach the final location. The length to any accepting location is exactly $j-i$ and each location has a single incoming transition.
2. Let $D = D_1 \frown D_2$. In this case $\mathcal{A}_D^{[i,j]}$ is given by $\bigcup_k [\mathcal{A}_{D_1}^{[i,k]} \cdot \mathcal{A}_{D_2}^{[k,j]}]$ for $i \leq k \leq j$. Here $\mathcal{A}_{D_1}^{[i,k]} \cdot \mathcal{A}_{D_2}^{[k,j]}$ fuses the final state of $\mathcal{A}_{D_1}^{[i,k]}$ with the initial location of $\mathcal{A}_{D_2}^{[k,j]}$, and the idea is to obtain the union of such concatenations for all $k \in [i, j]$. The concatenation of all behaviours of length $k-i$ satisfying D_1 over an interval I_1 followed by all possible behaviours of length $j-k$ satisfying D_2 over an interval I_2 gives all behaviours of length $j-i$ satisfying $D_1 \frown D_2$ over an interval I obtained as the fusion of intervals I_1, I_2 .

The constructed $\mathcal{A}_D^{[i,j]}$ is (i) acyclic, (ii) each location has at most one incoming transition, and (iii) the length of any path in $\mathcal{A}_D^{[i,j]}$ is $j-i$. Our goal is to obtain $\mathcal{A}_D^{[0,n]}$ for some bound n . This final automaton is obtained by taking the product of component automata, or fusing them as above,




Figure 2: A stopwatch automaton over the alphabet $\Sigma = 2^{Var}$ where $Var = \{p, q\}$. We have a stopwatch variable x_p and a clock y . The rate of x_p is 0 at locations q_0, q_2, q_4, q_5 . It can be seen that the accumulated duration of time for which p is true is some $t > 2$ on reaching q_2, q_3 via q_1 . Thereafter, on reaching q_5 , this duration is in $(t + 1, t + 2)$ for $t > 2$. Note that no duration is accumulated in location q_2 where 4 units of time are spent.

or in some cases, taking a complement. Complementation is not a problem here since we deal only with acyclic automata and bounded length behaviours.

4.2 Emptiness Problem

Theorem 3. *Acyclic stopwatch (timed) automata as obtained above have a decidable reachability.*

Consider an acyclic stopwatch automata \mathcal{A} as above, consisting of k locations s_1, \dots, s_k . Each of the $k - 1$ edges in any path is decorated by some Var_P . We encode the transitions and paths of $\mathcal{A}_D^{[i,j]}$ as QF_LRA formulae¹. Being acyclic, the whole automaton is then encoded as a disjunction of finitely many QF_LRA formulae, depending on the number of branches/paths the automaton has : each path leading to an accepting state is encoded by a QF_LRA formula. The idea is that $\mathcal{A}_D^{[i,j]}$ has a path leading to an accepting state if we can satisfy the QF_LRA formula.

The linear constraints we consider in our QF_LRA formulae are of the form $x - y \bowtie c, \sum_i x_i \bowtie c$ for real variables x_i, x, y and $c \in \mathbb{N}$. The constraints of the form $x - y \bowtie c$ are useful in encoding the difference between real valued variables representing time elapses while $\sum_i x_i \bowtie c$ is useful in encoding the accumulated duration of a proposition P across several states (variable x_i encodes time elapse in a state i). The Boolean combinations of propositional variables allowed in our QF_LRA formulae are handy to encode the labels on the transitions.

Checking the satisfaction of the QF_LRA formula amounts to checking if we have an interpretation which satisfies all the constraints. Substituting a fresh propositional variable for each distinct constraint in this QF_LRA formula, we obtain a propositional logic formula ζ over an extended set of variables $Var' \supset Var$. If we obtain a satisfying assignment for the variables in Var' , which satisfies this formula, then we also know that the original QF_LRA formula is satisfiable. Guessing such an assignment non-deterministically, we can verify if it forms a satisfying assignment or not. The size of ζ is linear in the size of our QF_LRA formula, hence checking if the guessed assignment

¹A QF_LRA formula is a quantifier-free Boolean combination of propositional variables and linear constraints over real variables.

satisfies ζ takes time linear in the size of the QF_LRA formula. This way we get a non-deterministic polynomial time procedure (in NP) to check satisfiability of the QF_LRA formula, and hence the non emptiness of our automaton also. The lower bound to show NP-hardness is easy, since any instance of a SAT problem is one of QF_LRA, and hence also for the paths of our automaton as well.

Theorem 4. *DC under the bounded semantics has a decidable satisfiability.*

Starting from a formula in DC under the bounded semantics, we first construct the timed or stopwatch automaton corresponding to it as described above. This construction is such that any bounded timed trace satisfying the formula is accepted by the constructed automata. The decidability of emptiness checking in the constructed automata is shown using Theorem 5, and is NP-complete.

4.3 Safety Games

When DC specifications concern the choices of two agents (the controller and the environment), the set of signals can be partitioned into controllable and uncontrollable signals. The corresponding translation to stopwatch automata gives rise to two-player safety game (or minimax reachability games) on stopwatch automata. In order to compute the maximally-permissive set of actions of a pacemaker, we need to solve a safety game on the resulting stopwatch automata. Similar to Theorem 4, we get the decidability of the reachability games

Theorem 5. *Safety games on acyclic stopwatch (timed) automata can be solved effectively.*

This safety region can be computed using controllable-predecessor operator and defines a set of maximally permissive actions for the pacemaker. When the underlying system model is timed automata, the winning region can be computed using UPPAAL-Tiga (Behrmann et al. 2007), and for a stopwatch automata it can be computed using Phaver+ (Benerecetti, Faella, and Minopoli 2011).

Once the safety region for the DC requirement has been computed, it can be used to block unsafe actions of the RL agent. This is the crux of shielding based reinforcement learning (Alshiekh et al. 2018).

5 Experimental Results

We have implemented DC to stopwatch automata construction in an extension of DCVALID (dcvalid) with some optimizations and derived operators. DCVALID in turn uses MONA (Klarlund and Møller 2001) and provides a validity checker and visualizer for DC Formulae.

5.1 Validation of Pacemaker Requirements

We have written the specification of pacemaker in DC logic. We present the full list of requirements in Appendix B. We used DCVALID to compile the pacemaker specification into automata and validated their correctness using two requirements shown in Section A.2 and A.3 (in Appendix). In order to simulate the pacemaker, we compiled each requirement into an automaton, resulting in a large number of implicitly




Figure 3: Ventricular Safety Pacing




Figure 4: Upper Rate Holdoff

conjoined automata. Since explicitly computing their product is neither practical (due to the state space blowup and the resulting need for minimization) nor necessary for the simulation, we implemented an on-the-fly simulation algorithm.

The simulation algorithm begins by determining a topology of the automata through ordering them with respect to their inter-dependency. The inputs to the system are provided by simulating the intrinsic atrial and ventricular heart beats, and each automaton is simulated in the order determined by the topology, which in turn generates the output events of atrial and ventricular pacing. Here are the results of our simulation on two scenarios (Figures 8 and 9).

Test Case 1 (Ventricular Safety Pacing). Atrial paced events are followed by a Ventricular Safety Pace (VSP) period (shown as signal *SAFE* in Figure 3). Any VS detected while this signal is asserted will trigger a VP when the signal ends. As seen in the figure, *safeVP* indeed occurs on the falling edge of *SAFE* showing the VSP functionality.

Test Case 2 (Upper Rate Hold-off). The scenario depicted in Figure 9 is a slightly complex: if VP is triggered at the precise time, the upper and lower rate requirements could be violated. The presented test case in Figure 4 shows satisfaction of this complicated situation.

Since the first ventricular sense in the example occurs within the refractory period, it can be ignored and thus it does not change any pacing timing. Meanwhile, the subsequent atrial sensed event, AS, occurs quite early. If the ventricular sensed event was not ignored, a new VA interval would have started causing the AS to fall into refractory interval PVARP. This would have ignored the AS setting the subsequent VP to occur at LRL. However, as seen in Figure 4, *PVAR* does not get restarted after the refractory VS (the first pulse on the *VS* trace.) Additionally, the AS did cause assertion of the *SAV* signal showing the start of a new AV interval. This too would not have occurred if the refractory VS had been misclassified. Finally, because of the early arrival of the AS, the upper rate *URL* signal is still asserted and thus, the scheduled VP must be held off till the URL ends. Once *URL* completes, signals *lateVP* is asserted showing correct hold-off.

5.2 Design of an Adaptive DDD Pacemaker

Consider the Mobitz II second-degree AV discussed in Example 1. Figure 5a shows a short timeline of 3:2 heart block pacing. The pacemaker does not look for the underlying intrinsic rate and will pace at the end of the lower rate interval. Since the patients can feel this difference in heartbeats, we design an adaptive pacemaker that finds an optimal time to provide for a missing beat matching the intrinsic rhythm.

For this case-study, we start from the DC specification of the DDD pacemaker presented in Appendix B with one exception. In a standard pacemaker, the timing of the ventricular pace is fixed in time, occurring at the termination of the AV interval. For this example, this requirement was relaxed, allowing the pace to occur any time in the AV interval. The modified pacemaker specification is presented in Appendix B.4. By embodying all the pacemaker requirements, the automaton guards against incorrect RL actions, rejecting those that would violate a non-permissive requirement. We designed an RL agent with the information of the safety region of the resulting specification, thus eliminating unsafe choices from the RL agent. The RL agent was rewarded based on how closely it matched the intrinsic rhythm of the heart model. To generalize the learning to different environments, we extended the state information with the history of the last 10 AV intervals.

For this case-study, we created a simple heart model with a 3:2 heart block. In addition, the heart model was permitted to randomly accelerate and decelerate with a 10% chance in either direction. Figure 6 summarizes the timing of the intrinsic and paced events over a run of 286 cardiac cycles. The light green line represents each atrial contractions. The dark green and red are the total ventricular intrinsic beats and the pacemaker induced respectively. The pacemaker was set to pace at an interval of 15 in lieu of any intrinsic ventricular activity. The RL agent learned that the optimal time to pace was two time periods after the missing intrinsic beat and on the chart each green bar is paired with a red bar approximately half its size and to the right. It is desired to always allow the heart to intrinsically beat if possible, thus, two time periods of delay was enforced by the reward generator to assure the pacemaker does not begin to pace over the heart.

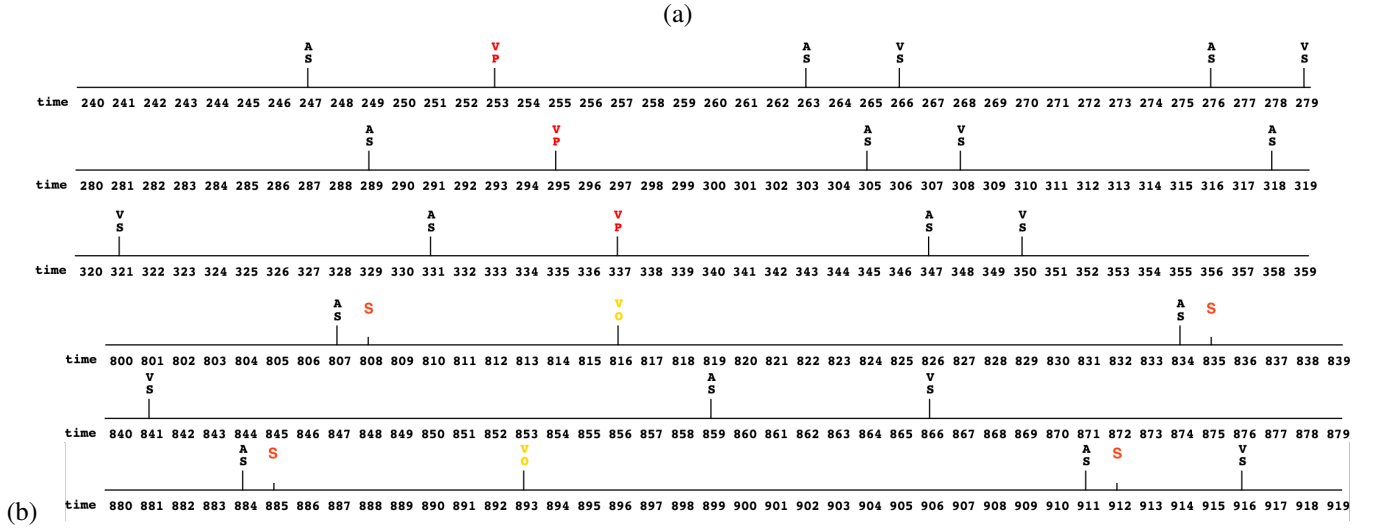


Figure 5: (a) Non-permissive pacemaker synthesized via RL pacing heart with 3:2 heart block; (b) An adaptive pacemaker learned via RL, shield from safety specifications, heart with 3:2 heart block.

The pace totals in the figure is not exact as the RL agent undergoes an adjustment period each time the intrinsic rate changed. The optimal time for an intrinsic beat at 7 here is at 10 because pacing at 9 would violate the pacemaker’s URL requirement so it had to be held off.

Figure 5 shows part of a timeline of the experiment run. An example of the optimal timing can be seen by looking at a sequence starting at time 834 with an atrial sense. The ventricle intrinsically beats at 841. This pattern repeats at times 859 and 866. The next atrial event at 884 does not have an intrinsic beat at 891. This is the dropped beat. The RL agent provides support at 893 or two increments later as expected. This timeline also demonstrate the utility of the safety shield in action. In many cases, atrial senses are followed immediately by the agent attempting to pace. This time is the blanking period and is not permitted by the pacemaker requirements. These actions, denoted by an ‘S’ symbol, were prevented from occurring by the shield.

5.3 Discussions

We looked at the generalizing nature of RL to create new and unique features in medical devices and other real-time control systems. We presented an improvement in patient comfort by regulating pacing to match the intrinsic heart rate of heart block patients. A major concern while deploying RL in safety critical systems is that allowing unfettered exploration to find optimal paths may exceed safety parameters. Shielding the exploration by limiting the RL’s action choice to remain in safety zones allows the RL agent to safely train. Once trained, the agent can be deployed with no further learning while still constrained by the shield to prevent any possible future erroneous actions. When constrained to specific problems, RL agents trained using locally permissive requirements guarded from exceeding safety parameters can provide a way to add new features to safety critical systems.

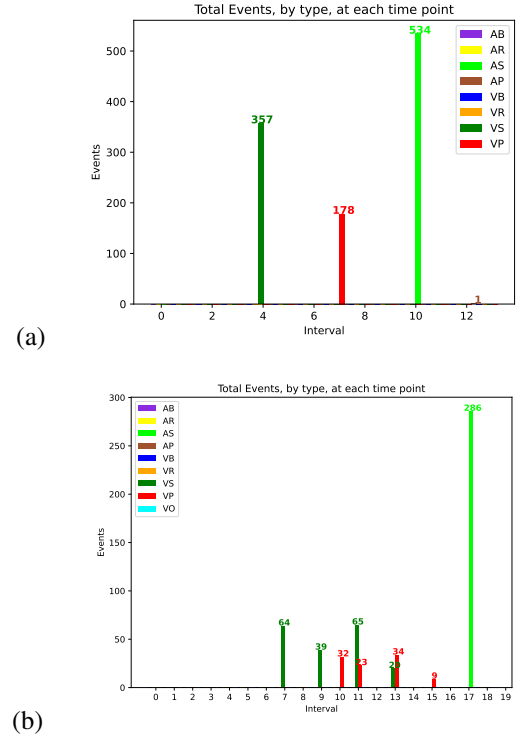


Figure 6: Timing with a) Standard Pacemaker and b) Permissive Pacemaker; 3:2 heart block

References

- Alshiekh, M.; Bloem, R.; Ehlers, R.; Könighofer, B.; Niekum, S.; and Topcu, U. 2018. Safe reinforcement learning via shielding. In *Thirty-Second AAAI Conference on Artificial Intelligence*.
- Alur, R.; and Dill, D. L. 1994. A Theory of Timed Automata.

Theoretical Computer Science, 126(2): 183–235.

Baier, C.; and Katoen, J. 2008. *Principles of model checking*. MIT Press.

Behrmann, G.; Cougnard, A.; David, A.; Fleury, E.; Larsen, K. G.; and Lime, D. 2007. UPPAAL-Tiga: Time for Playing Games! In *Proceedings of the 19th International Conference on Computer Aided Verification*, 121–125.

Benerecetti, M.; Faella, M.; and Minopoli, S. 2011. Automatic Synthesis of Switching Controllers for Linear Hybrid Automata.

Bernstein, A. D.; Daubert, J.-C.; Fletcher, R. D.; Hayes, D. L.; Lüderitz, B.; Reynolds, D. W.; Schoenfeld, M. H.; and Sutton, R. 2002. The revised NASPE/BPEG generic code for atribradycardia, adaptive-rate, and multisite pacing. *Pacing and clinical electrophysiology*, 25(2): 260–264.

Bonfanti, S.; Gargantini, A.; and Mashkoor, A. 2018. A systematic literature review of the use of formal methods in medical software systems. *Journal of Software: Evolution and Process*, 30(5): e1943.

Cassez, F.; and Larsen, K. 2000. The Impressive Power of Stopwatches. In Palamidessi, C., ed., *CONCUR 2000 — Concurrency Theory*, 138–152. Berlin, Heidelberg: Springer Berlin Heidelberg.

Chaochen, Z.; Hansen, M. R.; and Sestoft, P. 1993. Decidability and Undecidability Results for Duration Calculus. In *STACS 93, 10th Annual Symposium on Theoretical Aspects of Computer Science, Würzburg, Germany, February 25-27, 1993, Proceedings*, 58–68.

Chaochen, Z.; Hoare, C. A. R.; and Ravn, A. P. 1991. A Calculus of Durations. *Inf. Process. Lett.*, 40(5): 269–276.

devalid. ??? DCVALID—A tool for modelchecking Duration Calculus Formulae. accessed: 05/28/2021.

Dole, K.; Gupta, A.; and Krishna, S. N. 2020. Robust Controller Synthesis for Duration Calculus. In *International Symposium on Automated Technology for Verification and Analysis*, 429–446. Springer.

Finkelmeier, N. E. 1991. Pacemaker technology: an overview. *AACN Advanced Critical Care*, 2(1): 99–106.

Gomes, A. O.; and Oliveira, M. V. M. 2009. Formal Specification of a Cardiac Pacing System. In Cavalcanti, A.; and Dams, D. R., eds., *FM 2009: Formal Methods*, 692–707. Berlin, Heidelberg: Springer Berlin Heidelberg.

Gomes, J. A. 2020. The Artificial Pacemaker: A Historical Overview. In *Heart Rhythm Disorders*, 449–460. Springer.

Greenspon, A. J.; Patel, J. D.; Lau, E.; Ochoa, J. A.; Frisch, D. R.; Ho, R. T.; Pavri, B. B.; and Kurtz, S. M. 2012. Trends in permanent pacemaker implantation in the United States from 1993 to 2009: increasing complexity of patients and procedures. *Journal of the American College of Cardiology*, 60(16): 1540–1545.

Henzinger, T. A.; Kopke, P. W.; Puri, A.; and Varaiya, P. 1995. What’s decidable about hybrid automata? In *Proceedings of the Twenty-Seventh Annual ACM Symposium on Theory of Computing*, 29 May-1 June 1995, Las Vegas, Nevada, USA, 373–382.

Jiang, Z.; Pajic, M.; Moarref, S.; Alur, R.; and Mangharam, R. 2012. Modeling and verification of a dual chamber implantable pacemaker. In *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*, 188–203. Springer.

Kay, G. N.; and Shepard, R. B. 2017. Stimulation and excitation of cardiac tissues. In *Clinical Cardiac Pacing, Defibrillation and Resynchronization Therapy*, 61–113. Elsevier.

Klarlund, N.; and Møller, A. 2001. *MONA Version 1.4 User Manual*. BRICS, Department of Computer Science, University of Aarhus. Notes Series NS-01-1. Available from <http://www.brics.dk/mona/>.

Könighofer, B.; Lorber, F.; Jansen, N.; and Bloem, R. 2020. Shield synthesis for reinforcement learning. In *International Symposium on Leveraging Applications of Formal Methods*, 290–306. Springer.

Laboratory, S. Q. R. 2007. Pacemaker System Specification. Langendorf, R.; and Pick, A. 1968. Atrioventricular block, type II (Mobitz)—its nature and clinical significance. *Circulation*, 38(5): 819–821.

Larson, B. R. 2014. Formal semantics for the pacemaker system specification. *ACM SIGAda Ada Letters*, 34(3): 47–60.

Lau, C.-P.; Siu, C.-W.; and Tse, H.-F. 2017. Sensors for Implantable Cardiac Pacing Devices. *Clinical Cardiac Pacing, Defibrillation and Resynchronization Therapy*, 281–312.

Maisel, W. H.; Sweeney, M. O.; Stevenson, W. G.; Ellison, K. E.; and Epstein, L. M. 2001. Recalls and safety alerts involving pacemakers and implantable cardioverter-defibrillator generators. *Jama*, 286(7): 793–799.

Méry, D.; Schätz, B.; and Wasssyng, A. 2014. The pacemaker challenge: developing certifiable medical devices (dagstuhl seminar 14062). In *Dagstuhl Reports*, volume 4.2. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik.

Ouaknine, J.; Rabinovich, A.; and Worrell, J. 2009. Time-Bounded Verification. In *CONCUR 2009 - Concurrency Theory, 20th International Conference, CONCUR 2009, Bologna, Italy, September 1-4, 2009. Proceedings*, 496–510.

Pandya, P. K. 2001. Specifying and Deciding Quantified Discrete-time Duration Calculus Formulae using DC-VALID: An Automata Theoretic Approach. In *Proceedings of RTTOOLS*.

Pandya, P. K. 2002. Interval Duration Logic: Expressiveness and Decidability. *Electron. Notes Theor. Comput. Sci.*, 65(6): 254–272.

Polyzos, K. A.; Konstantelias, A. A.; and Falagas, M. E. 2015. Risk factors for cardiac implantable electronic device infection: a systematic review and meta-analysis. *Ep Europace*, 17(5): 767–777.

Sutton, R. S.; and Barto, A. G. 2018. *Reinforcement learning: An introduction*. MIT press.

Vijayaraman, P.; and Ellenbogen, K. A. 2017. Atrioventricular Conduction System Disease. In *Clinical Cardiac Pacing, Defibrillation and Resynchronization Therapy*, 399–453. Elsevier.

Technical Appendix

Here we present more details needed for the technical appendix.

A Real-Time Requirements of Pacemakers

An electrocardiogram (ECG) is a noninvasive procedure that measures and records the electrical activity (voltage versus time) of the heart. This section introduces a graphical representation for the heart-pacemaker closed-loop system using a combined ECG graph and a so-called pacemaker timing ladder diagram. After explaining how to read these diagrams in the first subsection, we showcase two crucial scenarios of the pacemaker. We validate the DC specifications for the DDD pacemaker using these requirements in Section 5.1.

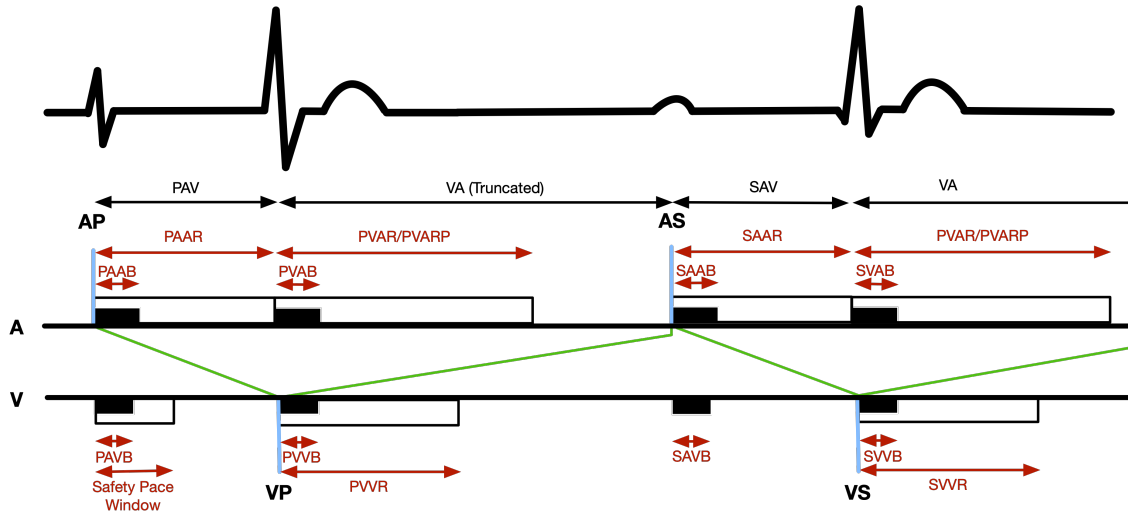


Figure 7: An ECG (top) and pacemaker timing ladder diagram (bottom) showing two cardiac cycles. The first cycle represents a paced atrial and ventricular events (AP,VP) in the ladder diagram. The second cycle represents intrinsic heart rhythm with atrial and ventricular sensed events (AS,VS). The post VP VA interval has been truncated by the early AS arrival.

A.1 ECG and the Timing Ladder Diagram

The waveform at the top of Figure 7 is a simple representation of heart activity as seen in a surface ECG. The waveform begins with the pacemaker delivering pacing pulses to the atrium and ventricle. Pacing activity is recognized by the sharp signal spikes that extend well above and below baseline followed by the T-wave, a large hump in the ECG waveform. This is followed by intrinsic paces by the heart of the two chambers, again followed by the ventricular refractory T-wave. The ladder diagram at the bottom of Figure 7 depicts how a pacemaker views the timing relations between cardiac activity.

The AV and VA intervals. The top and bottom horizontal black lines represent timing activities in the atrium and ventricle, respectively. Cardiac events are represented by vertical lines and labeled with a two letter acronym where the first letter represents the event chamber: A(trium) or V(entricle) and the second is the event: S(ense) or P(ace). An atrial sense would have the label AS. The cardiac cycle is measured from atrial event to atrial event and is broken into two halves: *the AV interval* (the time from the atrial event to the ventricular event) and *the VA interval* (the remainder of the cardiac cycle from the ventricular event to the next atrial event). The AV interval can be further identified as a sensed (SAV) or paced (PAV) interval based on the event that triggered it. These intervals are displayed at the top of the ladder diagram and are synchronized with ECG.

Blanking and Refractory Periods. Each interval has four timing periods that start coincident with the event that began the interval. The *blanking periods* provide noise rejection immediately after a cardiac event and all cardiac activity is ignored during this time. *Refractory periods* represent the time required for the heart chamber to recover from the event. While these are considered valid cardiac signals, they occur too soon in the cardiac cycle. These periods are given a four letter acronym where the first letter denotes the event type (S/P), the second is the chamber of the event (A/V), the third is the chamber of the timer (A/V), and the fourth is the type of the period (B/R).

The atrial blanking period following a sensed ventricular event would be SVAB. Following a ventricular event, an atrial blanking period (PVAB/SVAB) and refractory period (PVAR/SVAR) will begin along with similar periods in the ventricle (PVVB/PVVR). Similar periods occur after an atrial event. There are a few exceptions. By tradition, atrial sensed and paced refractory periods following a ventricular event are referred to as PVAR. There is no ventricular refractory period after an atrial




Figure 8: Safety Pacing and Atrial Holdoff




Figure 9: Upper Rate Holdoff and Ventricular Refractory Sensed Event

event. Instead, a safety pace period may exist. If a ventricular sensed event occurs during the safety pace window, a back up or safety pace is generated at the end of the window to provide additional hemodynamic support (see Example A.2).

The green line in the diagram shows the progressing of the interval time. It starts at the event on the timeline of the chamber where it occurred. If it reaches the opposite timeline without an intrinsic cardiac event, a pace in the chamber associated with the reached timeline is generated. The following scenarios graphically explain two key requirements.

A.2 Safety Pacing and Atrial Holdoff

Requirement. *In absence of intrinsic atrial activity, the pacemaker attempts to maintain the time between atrial beats at the Lower Rate Limit (LRL).*

Example. In Figure 8, the lower rate is 60 bpm (1000ms). A ventricular sensed event was detected at 75ms after the atrial pace. This VS falls into the safety pace period. The PAV interval continues as indicated by the green line until the safety pace window closes when a ventricular backup pace occurs 150ms after the atrial pace and the VA interval begins. At the end of the VA interval, the next atrial pace must be held off. A pace at the completion of the VA interval would occur 100ms too soon because of the back pace shortened PAV interval and would result in a pacing rate faster than lower limit. The atrial pace ultimately occurs 1000ms after the last atrial pace maintaining the A-A timing of the LRL.

A.3 Upper Rate Holdoff and Ventricular Refractory Sense

Requirement. *A pacemaker must never pace above the Upper Rate Limit (URL).*

Example. In Figure 9, the example shows two conditions that needs to be address by the pacemaker. Here, the cardiac cycle begins with complete pacing support with an AP followed by a VP at the conclusion of the PAV interval (200ms). A sensed ventricular event occurs at 250ms. The PVVR period is still active so this event is labeled a refractory ventricular sense (VR). The PVVR timer is restarted to allow for ventricular repolarization to complete after the VR but the overall VA interval is not




Figure 10: A visual explanation of the derived operators.

affected. An atrial sense occurs at 350ms after the VP. Being outside of PVAR, this is a valid atrial event and the AV interval begins. At the completion of the AV interval, 550ms have passed since the VP while the URL is 600ms. The VP must be held off for 50ms to prevent violation of the URL requirement. The VA interval still starts at the end of the AV interval to maintain the LRL rate of 1000ms. If the AV interval was also held off until the VP, the next atrial pace would occur at 1050ms, a violation of LRL.

B DC Specifications for Pacemakers

Now we are in a position to present the DC specifications for the DDD pacemaker requirements from (Laboratory 2007). We start this section by introducing some derived operators of DC that allow us to succinctly express various idioms in pacemaker specifications. In Section B.2 we introduce various signals and terms used in our specifications, while in Section B.3, we use these operators and signals to express DDD requirements.

B.1 Derived Operators: Syntactic Sugar for DC

We present some derived DC operators that allow us to express our specifications more succinctly.

1. We extend the usage of $\lceil X \rceil^\bullet$ and $\lceil X \rceil$ over Boolean combination of predicates as follows.

$$\lceil X \otimes Y \rceil^\bullet \stackrel{\text{def}}{=} \lceil X \rceil^\bullet \otimes \lceil Y \rceil^\bullet, \quad \lceil \neg X \rceil^\bullet \stackrel{\text{def}}{=} \neg \lceil X \rceil^\bullet, \quad \lceil X \rceil \stackrel{\text{def}}{=} \Box(\lceil X \rceil^\bullet \frown \text{true}), \quad \text{and} \quad \lceil X \rceil^d \stackrel{\text{def}}{=} (\ell = d \wedge \lceil X \rceil) \frown \text{true}$$

where \otimes is some logical binary operator and X and Y are Boolean formulae over predicates.

2. The operator $A \hookrightarrow_d B$ encodes that A triggers B for precisely d time units if A occurs again before d time then B is re-triggers, i.e.,

$$A \hookrightarrow_d B \stackrel{\text{def}}{=} \Box \left((\ell \geq d + 1) \wedge (\lceil A \rceil^\bullet \frown \text{true}) \Rightarrow (\ell = 1) \frown \lceil B \rceil^d \right) \wedge \Box \left((\ell \geq d + 1) \wedge \lceil \neg A \rceil^d \Rightarrow \ell = d + 1 \frown \lceil \neg B \rceil^\bullet \frown \text{true} \right)$$

3. The operator $A \xleftrightarrow{C} B$ encodes that C holds between events A and B , i.e.,

$$A \xleftrightarrow{C} B \stackrel{\text{def}}{=} \Box(\ell \geq 1 \wedge (\lceil A \wedge \neg B \rceil^\bullet \frown \text{true}) \Rightarrow (\ell = 1 \frown \text{waitFor}(B, C))) \wedge \Box(\ell \geq 1 \wedge (\lceil B \rceil^\bullet \frown \text{true}) \Rightarrow (\ell = 1 \frown \text{waitFor}(A, \neg C)))$$

where *waitFor* is defined as follows.

$$\text{waitFor}(X, Y) \stackrel{\text{def}}{=} \lceil \neg X \wedge Y \rceil \vee (\lceil \neg X \wedge Y \rceil \frown \ell = 1 \frown \lceil X \wedge Y \rceil^\bullet \frown \text{true}) \vee (\lceil X \wedge Y \rceil^\bullet \frown \text{true})$$

4. The operator $A \frown B$ encodes that B is true when A transitions from true to false, i.e.,

$$A \frown B \stackrel{\text{def}}{=} \Box(\lceil A \rceil^\bullet \frown (\ell = 1) \frown \lceil \neg A \rceil^\bullet \Leftrightarrow (\ell = 1) \frown \lceil B \rceil^\bullet) \frown \text{true}.$$

5. Operator $A \xrightarrow{B?}_d C$ encodes C will be true after A and turns false after either a delay of time d or arrival of B ,

$$A \xrightarrow{B?}_d C \stackrel{\text{def}}{=} (A \hookrightarrow_d X) \wedge (A \xleftrightarrow{Y} B) \wedge \lceil (X \wedge Y) \Leftrightarrow C \rceil$$

where X and Y are intermediate signals.

6. Operator $(A \wedge B) \xrightarrow{\neg B} C$ encodes that C will be true when B turns false after both A and B are true,

$$(A \wedge B) \xrightarrow{\neg B} C \stackrel{\text{def}}{=} (B \curvearrowright X) \wedge (A \wedge B \xrightarrow{Y} X) \wedge (Y \curvearrowright C)$$

where X and Y are intermediate signals.

B.2 Signals, Time Intervals, and Derived Signals

Our specifications interpret DC formulae over signals therefore we express the specifications with the help of signals. In this modelling, we follow the convention that signal variable names are given in capital letters. A pacemaker receives the “sensing events” as input, while outputs the “pacing events.” Depending on the real-time constraints enforced by the requirements, the pacemaker decides whether it is appropriate to produce a pace signal at any given time instant. The time constraints are encoded with the help of derived signals. We first describe the input and outputs signals and then present the derived signals with their intuitive interpretations.

Input and output signals. The VS (ventricular sense) and AS (atrial sense) signals denote the heart activity sensed in the ventricle and atrium respectively. The output pacing signals of the ventricle and atrium are VP and AP. Using the sensed signals as input, the pacemaker decides when to produce VP and AP. We express the conditions that either enable or inhibit the pacing signals using derived signals that depend on the input and output signals.

Time interval constants. The following time constants define various configurable intervals of a pacemaker. We write them in lowercase letters. We will use them to define derived signals and write specifications.

- The `lrl` (lower rate limit) interval is the longest allowed time between two consecutive paced or ventricular sensed events, while the `url` (upper rate limit) interval is the shortest time possible between two consecutive ventricular paced or sensed events.
- The `pav` (paced atrium-ventricle) and `sav` (sensed atrium-ventricle) intervals start from a paced atrial event and sensed atrial event, respectively.
- The `va` (ventricle-atrial) interval starts after each non-refractory VS or after the expiration of `pav` or `sav`.
- The `pvab`, `pavb`, `pvvb`, `paab`, `svab`, `savb`, `svvb` and `saab` (paced/sensed ventricular/atrial atrial/ventricular blanking) intervals start at paced/sensed events at their second letter and affect the events in the third letter.
- The `pvvr`, `pvar`, `paar`, `svvr`, `svar` and `saar` (paced/sensed ventricular/atrial atrial/ventricular refractory) intervals are defined in an analogous fashion to the previous item.
- `safe` interval is used in case of enabling of safety ventricular pacing.

Derived signals. We present the derived signals that we need to write the specifications. We have divided the derived signals depending on the signals and timing information used for the construction of the signals.

- **Refractory and blanking signals.** For each paced or sensed (P or S) events in each chamber (A or V), we need to keep track of refractory or blanking (R or B) intervals that may affect any of the chambers (again A or V). Using the four letters, we describe 16 signals in writing our specification. For example, signal `PVVR` is true for `pvvr` period after a paced ventricular event. We will not use two of the sixteen, namely `SAVR` and `PAVR`, since they have no biological significance. In some circumstances, we need extended `PVAR` for which we used signal `ExPVAR`.
- **Accepted/non-ignored and non-refractory sensed signals.** The sensed events are sometimes masked. To achieve modularity, we define four signals `VSA`, `ASA`, `VSN`, and `ASN`, when each sensed events (V or A) are not masked by blanking or refractory period (A or N).
- **AV/VA interval signals.** If sufficient time has been passed since last paced or sensed (P or S) atrial event then we schedule a ventricular pace event if no other condition inhibits. We use either of two signals `PAV` and `SAV` to indicate the state of the interval depending on the event triggering the interval. We use signal `VA` to indicate the similar timing requirement from ventricular event to atrial pace event. For the VA interval, the specification does not differentiate between paced and sensed events in starting the interval. We also define signals `EPAV`, `ESAV` and `EVA` to indicate the 1 to 0 transitions in `PAV`, `SAV` and `VA` respectively.
- **Pacing time delay signals.** We use `LRL` and `URL` signals to indicate the lower and upper limit of pacing delays.
- **Different ventricular pacing signals.** There are different types of ventricular pace signals `scheduledVP`, `lateVP` and `safeVP` that are triggered in different circumstances. These signals together contribute to the final output ventricular pace (VP) signal.
- **Some special signals.** The signal `BOOT` is used as a trigger to start the system. Signal `NoAS` represents the absence of atrial sense after a non-refractory ventricular sensed event until an `ASA` event is seen.
- **Signals for permissive pacing.** The signal `VO` is input to our system along with the `AS` and `VS` signal to allow pacing decision from the reinforcement learning agent. The signal `Vother` is used to allow permissive pacing using the reinforcement learning agents input to our system.

B.3 Pacemaker Specifications

We divide our specifications in the following two groups. The first group defines the signals and delays which will be used further in writing the specifications. In the second group, we use the defined signals to write specification of the pacemaker.

1. The BOOT signal marks the start of the pacemaker.

$$[BOOT]^\bullet \vee [BOOT]^\bullet \wedge \ell = 1 \wedge [\neg BOOT]$$

2. A ventricular blanking period shall start immediately after a ventricular/atrial paced or sensed event. The blanking intervals $pvvb, svvb, pavb, savb$ are relevant to ventricular events.

$$(VP \hookrightarrow_{pvvb} PVVB) \wedge (AP \hookrightarrow_{pavb} PAVB) \wedge (VSA \hookrightarrow_{svvb} SVVB) \wedge (ASA \hookrightarrow_{savb} SAVB)$$

3. During the post atrial event, ventricular blanking period (PAVB/SAVB), ventricular sensed events shall be ignored. During the post ventricular event, ventricular blanking period (PVVB/SVVB), ventricular sensed events shall be ignored.

$$[VS \wedge \neg PVVB \wedge \neg PAVB \wedge \neg SVVB \wedge \neg SAVB \Leftrightarrow VSA]$$

4. An atrial blanking period shall start immediately after a atrial/ventricular paced or sensed event. $paab, saab, pvab, svab$ are the blanking intervals relevant to atrial events.

$$(VP \hookrightarrow_{pvab} PVAB) \wedge (AP \hookrightarrow_{paab} PAAB) \wedge (VSA \hookrightarrow_{svab} SVAB) \wedge (ASA \hookrightarrow_{saab} SAAB)$$

5. During the post atrial event, atrial blanking period (PAAB/SAAB), atrial sensed events shall be ignored. During the post ventricular event atrial blanking period (PVAB/SVAB), atrial sensed events shall be ignored.

$$[AS \wedge \neg PVAB \wedge \neg PAAB \wedge \neg SVAB \wedge \neg SAAB \Leftrightarrow ASA]$$

6. A post ventricular, atrial refractory period (PVAR) and post ventricular, ventricular refractory period (PVVR) shall start immediately after a non-refractory ventricular sensed or paced event.

$$((VP \vee VSA) \hookrightarrow_{pvvr} PVVR) \wedge ((VP \vee VSN \vee (ASA \wedge PVAR)) \hookrightarrow_{pvar} PVAR)$$

7. A ventricular sensed event is non-refractory if it occurs when PVVR is not true.

$$[\neg PVVR \wedge VSA \Leftrightarrow VSN]$$

8. A post atrial, atrial refractory period (PAAR) shall start immediately after a non-refractory atrial sensed or paced event. Note that the PAAR/SAAR runs the entire length of the PAV/SAV interval.

$$(AP \vee ASN) \xrightarrow{PAAR} (VP \vee VSN)$$

9. Atrial refractory runs the entire AV interval.

$$[\neg PAAR \wedge \neg PVAR \wedge \neg ExPVAR \wedge ASA \Leftrightarrow ASN]$$

10. A fixed paced AV (PAV) delay shall be initiated by a paced atrial event. A fixed sensed AV (SAV) delay shall be initiated by a sensed non-refractory atrial event. EPAV and ESAV signals mark the end point of the AV interval.

$$[VSN \Leftrightarrow Vother]$$

$$(AP \xrightarrow{Vother?}_{pav} PAV) \wedge (ASN \xrightarrow{Vother?}_{sav} SAV) \wedge$$

$$((PAV \wedge \neg Vother) \curvearrowright EPAV) \wedge ((SAV \wedge \neg Vother) \curvearrowright ESAV)$$

Please note that we have introduced a signal $Vother$, which is synonym of VSN . We need this signal to be able to adapt our specification to support learning agent later in section B.4.

11. A fixed VA delay shall be initiated after each non-refractory sensed ventricular event or after completion of the AV (SAV/-PAV) interval, whichever occurs first.

$$(Vother \vee ESAV \vee EPAV \vee BOOT) \xrightarrow{ASN?}_{va} VA \wedge ((VA \wedge \neg ASN) \curvearrowright EVA)$$

12. A non-refractory ventricular sensed event followed by a second non-refractory ventricular event without an intervening atrial sensed event (a premature ventricular contraction or PVC) shall start an extended PVAR.

$$(VSN \xrightarrow{NoAS} ASA) \wedge ((NoAS \wedge VSN) \hookrightarrow_{expvar} ExPVAR)$$

13. The LRL and URL interval starts at a paced ventricular event or non-refractory ventricular sensed event.

$$((VP \vee VSN) \hookrightarrow_{\text{lrl}} LRL) \wedge ((VP \vee VSN) \hookrightarrow_{\text{url}} URL) \wedge ((AP \vee ASN) \hookrightarrow_{\text{lrl}} ALRL) \wedge ((AP \vee ASN) \hookrightarrow_{\text{url}} AURL)$$

Specifications AP1-AP4 are for generating AP events and VP1-VP4 are for generating VP events. OE1 and OE2 define the operation envelop of a pacemaker.

- **[AP1]** VRP extension beyond the VA interval shall cause the atrium to be paced at the completion of the VA interval (noise reversion). *This condition does not add any restriction. Therefore, we do not need a formula.*
- **[AP2]** Ventricular sensed events outside of SVVR/PVVR shall begin a new VA interval. *This condition does not add any restriction. The definition of VA already handles the extension.*
- **[AP3]** Atrial sensed events inside of PVAR shall not start the SAV interval nor inhibit the scheduled atrial pace. Atrial sensed events outside of PVAR shall begin an SAV interval. *This condition does not add any restriction. The definition of VA already handles the end of VA.*
- **[AP4]** An atrial paced event shall occur upon completion of VA.

$$[(EVA \wedge \neg ASN \wedge \neg AURL) \Leftrightarrow \text{scheduledAP}]$$

- **[AP5]** If VA periods ends and URL timing has not met, the AP will wait for URL to be over.

$$((EVA \wedge \neg ASN) \wedge AURL) \xrightarrow{\neg AURL} \text{lateAP}, \quad [(\text{scheduledAP} \vee \text{lateAP}) \Leftrightarrow AP]$$

- **[VP1]** A non-refractory ventricular sensed event during the AV interval *outside the safety period* shall begin a VA interval. *This condition does not add any restriction. The definition of PAV/SAV already handles the end of PAV/SAV.*
- **[VP2]** A ventricular paced event shall occur upon completion of AV interval unless the upper rate limit has not been met.

$$[((PAV \vee ESAV) \wedge \neg VSN \wedge \neg URL) \Leftrightarrow \text{scheduledVP}]$$

- **[VP3]** If AV periods ends and URL timing has not met, the VP will wait for URL to be over.

$$(((PAV \vee ESAV) \wedge \neg V_{\text{other}}) \wedge URL) \xrightarrow{\neg URL} \text{lateVP}$$

- **[VP4]** A non-refractory ventricular sensed event during PAV interval inside safety period shall cause the next scheduled ventricular pace to occur at safety pace time.

$$(AP \hookrightarrow_{\text{safe}} SAFE) \wedge (VSN \wedge SAFE) \xrightarrow{\neg SAFE} \text{safeVP}, \quad [(\text{scheduledVP} \vee \text{lateVP} \vee \text{safeVP}) \Leftrightarrow VP]$$

- **[OE1]** The time between paces shall not exceed the lower rate limit interval (LRL).

$$[VP \Rightarrow LRL]$$

- **[OE2]** The time between a ventricular event and the next ventricular pace shall not be less than the Upper Rate Limit (URL) interval.

$$[VP \Rightarrow \neg URL]$$

This completes the set of requirements for DDD pacemaker from (Laboratory 2007). ■

B.4 Permissive Specification

To support the inputs from the learning agent, we introduce another input signal VO that indicates the pacing decision of the learning agent. We allow learning agent to interrupt PAV or SAV , when $SAFE$ and URL signals are false. To accodate VO , we modify the definition of V_{other} .

$$[VSN \vee (VO \wedge \neg SAFE \wedge \neg URL) \Leftrightarrow V_{\text{other}}]$$

The specification VP2 ensures that if the PAV or SAV intervals are ended by VO then we produce a scheduledVP , i.e., the learning agent is allowed to pace.